

How to Install Kubernetes Cluster (kubeadm) on Ubuntu 22.04

Kubernetes is a powerful container orchestration platform used for automating the deployment, scaling, and management of containerized applications. The step-by-step process of installing Kubernetes on Ubuntu 22.04. This cluster configuration includes a master node and worker nodes.

Kubernetes Nodes

In a Kubernetes cluster, we have two distinct categories of nodes:

Master Nodes: These nodes play a crucial role in managing the control API calls for various components within the Kubernetes cluster. This includes overseeing pods, replication controllers, services, nodes, and more.

Worker Nodes: Worker nodes are responsible for providing runtime environments for containers. It's worth noting that a group of

container pods can extend across multiple worker nodes, ensuring optimal resource allocation and management.

Prerequisites

Before diving into the installation, ensure that your environment meets the following prerequisites:

- An Ubuntu 22.04 system.
- Privileged access to the system (root or sudo user).
- Minimum 2GB RAM or more.
- Minimum 2 CPU cores (or 2 vCPUs).
- 20 GB of free disk space (or more).

Step 1: Update and Upgrade Ubuntu (all nodes)

Open a terminal and execute the following commands:

- `sudo apt-get update`

Step 2: Disable Swap (all nodes)

To enhance Kubernetes performance, disable swap and set essential kernel parameters.

- `sudo swapoff -a`
- `sudo sed -i 's/^(\.*)$/#\1/g' /etc/fstab` (or)
- `sudo vi /etc/fstab` (Comment the swapfile line & save)

Step 3: Add Kernel Parameters (all nodes)

Load the required kernel modules on all nodes:

- `cat <<EOF | sudo tee /etc/modules-load.d/k8s.conf`
`overlay`
`br_netfilter`
`EOF`
- `sudo modprobe overlay`
- `sudo modprobe br_netfilter`

Verify that the `br_netfilter`, `overlay` modules are loaded by running the following commands:

- `lsmod | grep br_netfilter`
- `lsmod | grep overlay`

Configure the critical kernel parameters for Kubernetes using the following:

- `sudo tee /etc/sysctl.d/kubernetes.conf <<EOF`
`net.bridge.bridge-nf-call-ip6tables = 1`
`net.bridge.bridge-nf-call-iptables = 1`
`net.ipv4.ip_forward = 1`

EOF

Then, reload the changes:

- `sudo sysctl -system`

This command is configuring some sysctl kernel parameters that are required for Kubernetes networking to function properly. Here's what it's doing:

`sudo tee /etc/sysctl.d/kubernetes.conf<<EOF` — Creates a new sysctl config file called `/etc/sysctl.d/kubernetes.conf` and writes the following parameters into it:

`net.bridge.bridge-nf-call-ip6tables = 1`

`net.bridge.bridge-nf-call-iptables = 1`

`net.ipv4.ip_forward = 1`

The first two parameters enable bridged IPv4 and IPv6 traffic to be passed to iptables chains. This is required for Kubernetes networking

policies and traffic routing to work.

`net.ipv4.ip_forward = 1` enables IP forwarding in the kernel, which is required for packet routing between pods in Kubernetes.

`sudo sysctl --system` — Applies the `sysctl` parameters from the new `/etc/sysctl.d/kubernetes.conf` file to the running kernel. This enables the settings without requiring a reboot.

In summary, this command configures three key `sysctl` parameters needed for Kubernetes networking and traffic policies and loads them into the running kernel so they are active immediately. The `/etc/sysctl.d/kubernetes.conf` file will persist these settings across reboots as well.

Step 4: Install Containerd Runtime (all nodes)

- `Wget`

https://download.docker.com/linux/ubuntu/dists/jammy/pool/stable/amd64/containerd.io_1.6.28-1_amd64.deb

- `sudo dpkg -i containerd.io_1.6.28-1_amd64.deb`

Configure containerd to start using systemd as cgroup:

- containerd config default | sudo tee /etc/containerd/config.toml

>/dev/null 2>&1
- sudo sed -i 's/SystemdCgroup \= false/SystemdCgroup \= true/g' /etc/containerd/config.toml

Restart and enable the containerd service:

- sudo systemctl restart containerd
- sudo systemctl enable containerd

Step 5: Add Apt Repository for Kubernetes (all nodes)

Kubernetes packages are not available in the default Ubuntu 22.04 repositories. Add the Kubernetes repositories with the following commands:

- sudo apt-get install -y apt-transport-https ca-certificates curl gpg
- curl -fsSL https://pkgs.k8s.io/core:/stable:/v1.29/deb/Release.key | sudo gpg --dearmor -o /etc/apt/keyrings/kubernetes-apt-keyring.gpg
- echo 'deb [signed-by=/etc/apt/keyrings/kubernetes-apt-keyring.gpg] https://pkgs.k8s.io/core:/stable:/v1.29/deb/ /' | sudo tee /etc/apt/sources.list.d/kubernetes.list

Step 6: Install Kubectl, Kubeadm, and Kubelet (all nodes)

After adding the repositories, install essential Kubernetes components, including kubectl, kubelet, and kubeadm, on all nodes with the following commands:

- `sudo apt-get update`
- `sudo apt-get install -y kubelet kubeadm kubectl`
- `sudo apt-mark hold kubelet kubeadm kubectl`

Step 7: Initialize Kubernetes Cluster with Kubeadm (master node)

- `sudo kubeadm init`

After the initialization is complete make a note of the **kubeadm join** command for future reference.

Run the following commands on the master node:

- `mkdir -p $HOME/.kube`
- `sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config`
- `sudo chown $(id -u):$(id -g) $HOME/.kube/config`

Step 8: Add Worker Nodes to the Cluster (Master nodes)

- `kubeadm token create --print-join-command`

Token will get Generated Copy and run the command on worker Node.

Step :9 Install Kubernetes Network Plugin (master node)

To enable communication between pods in the cluster, you need a network plugin. Install the Calico network plugin with the following command from the master node:

- `curl https://raw.githubusercontent.com/projectcalico/calico/v3.25.0/manifests/calico.yaml -O`
- `kubectl apply -f calico.yaml`

NAME	READY	STATUS	RESTARTS	AGE
calico-kube-controllers-658d97c59c-qghsz	1/1	Running	0	3d19h
calico-node-92cwg	1/1	Running	0	3d19h
calico-node-lz6gz	1/1	Running	1 (3d1h ago)	3d19h
calico-node-zz2t2	1/1	Running	0	3d19h

Step 10: Verify the cluster and test (master node)

Finally, we want to verify whether our cluster is successfully created.

- `kubectl get nodes`
- `kubectl get pods -n kube-system`


```
$ kubectl get nodes
STATUS    ROLES          AGE      VERSION
Ready     control-plane  3d20h    v1.29.3
Ready     <none>         3d20h    v1.29.3
Ready     <none>         3d19h    v1.29.3
```

```
c
```